



Apache Kafka

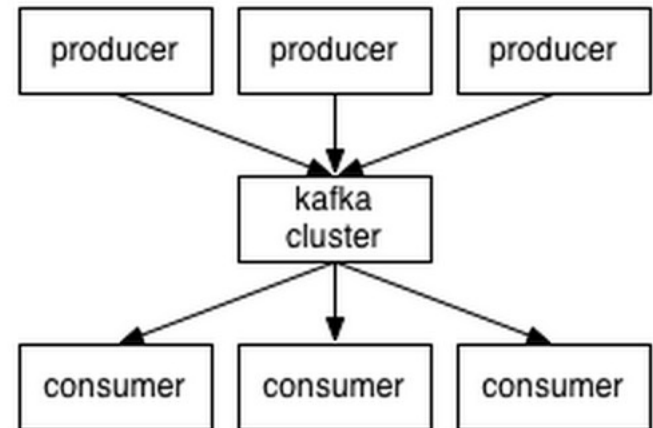
Slides modified from  
CMSC 491: Hadoop-Based Distributed Computing  
Spring 2016 Adam Shook

# Overview

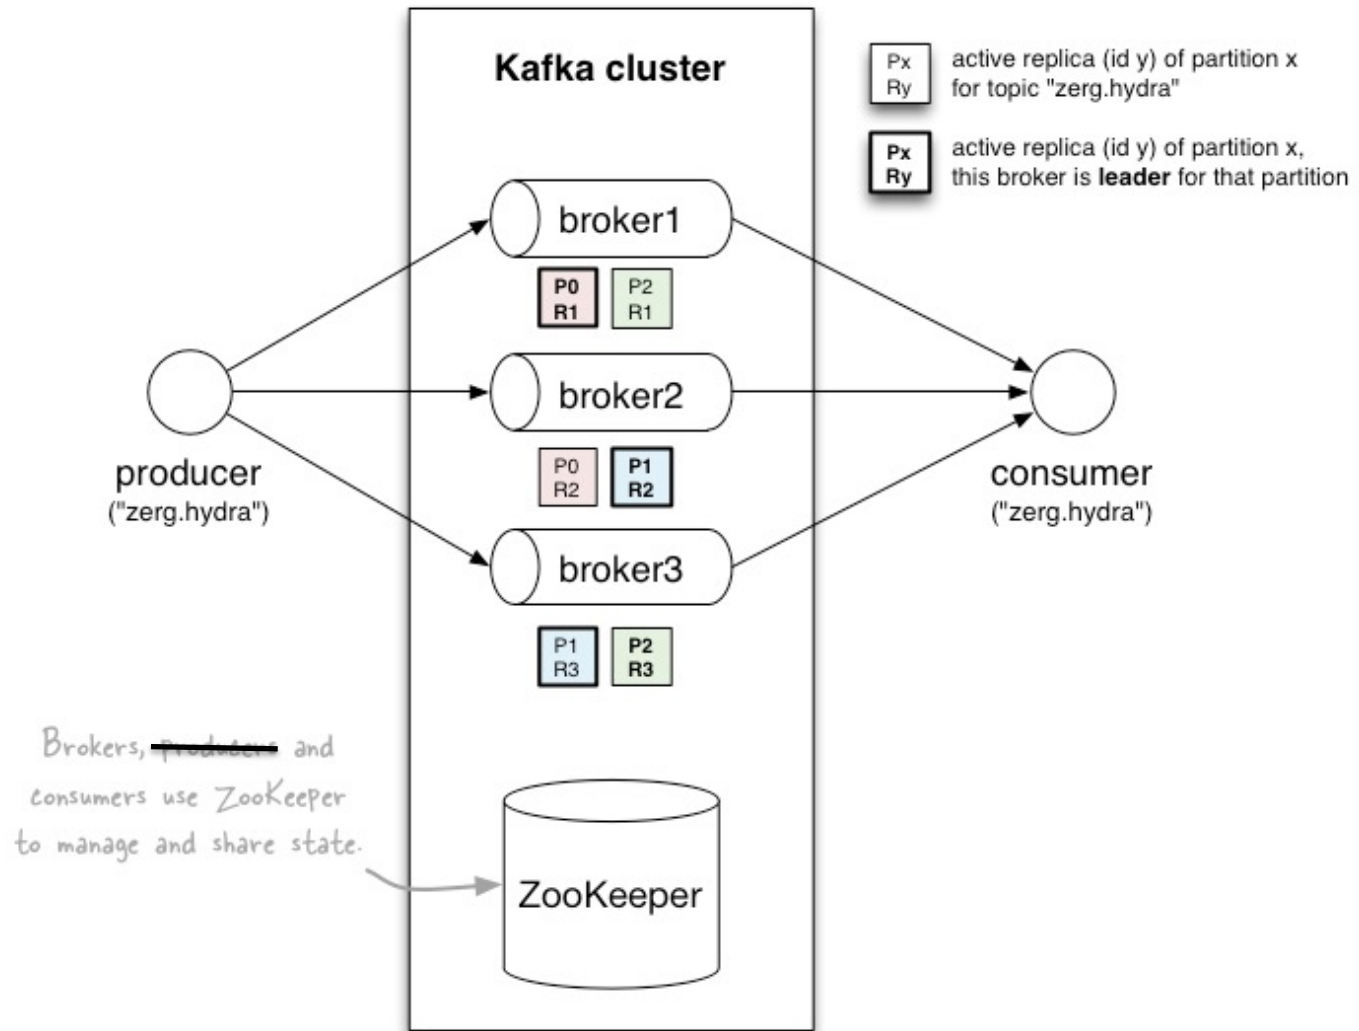
- Kafka is a “publish-subscribe messaging rethought as a distributed commit log”
- Fast
- Scalable
- Durable
- Distributed

# A first look

- The who is who
  - **Producers** write data to **brokers**.
  - **Consumers** read data from **brokers**.
  - All this is distributed, including
    - Producers
    - Kafka brokers and storage
    - Consumers
- The data
  - Data is stored in **topics**.
  - **Topics** are split into **partitions**, which are **replicated**.
  - Brokers run on different servers and manage data on the servers



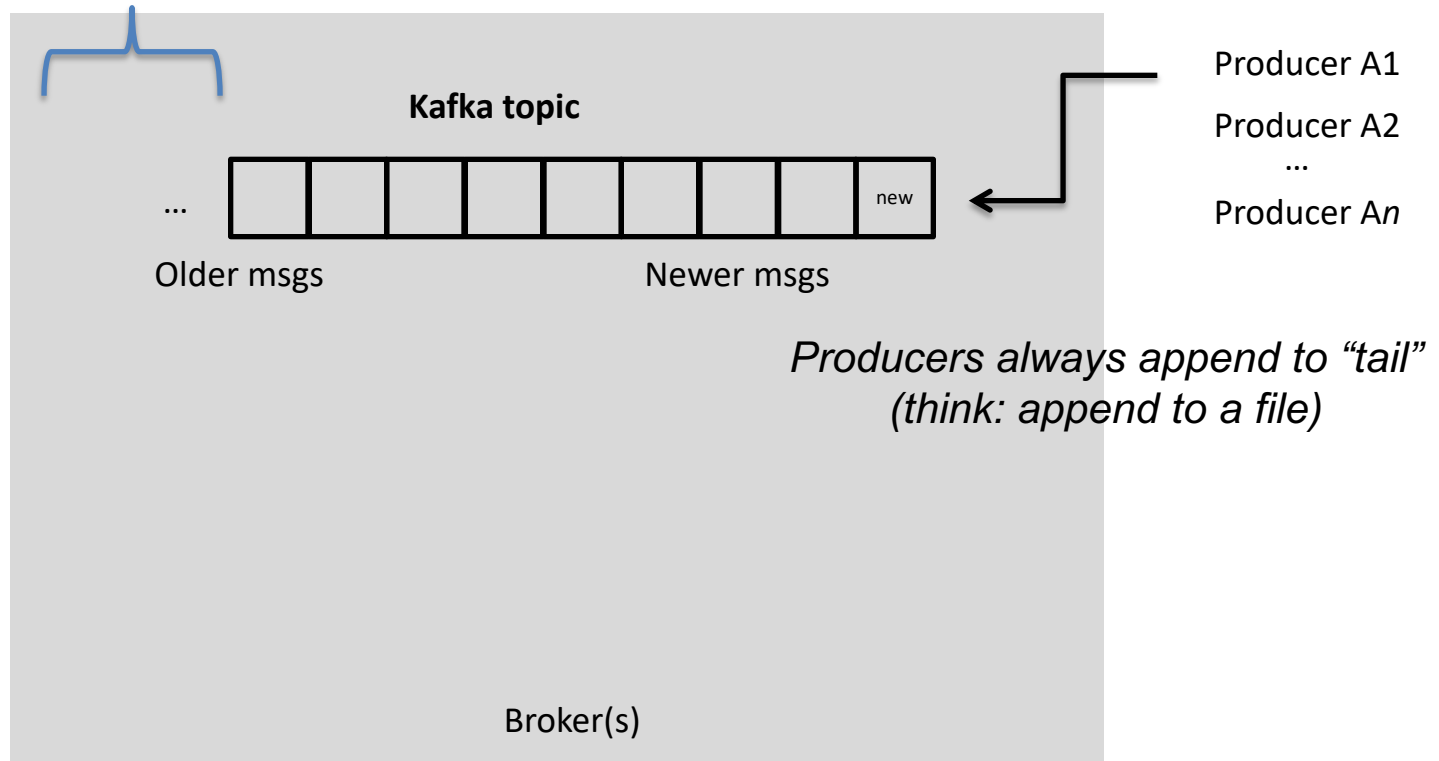
# A first look



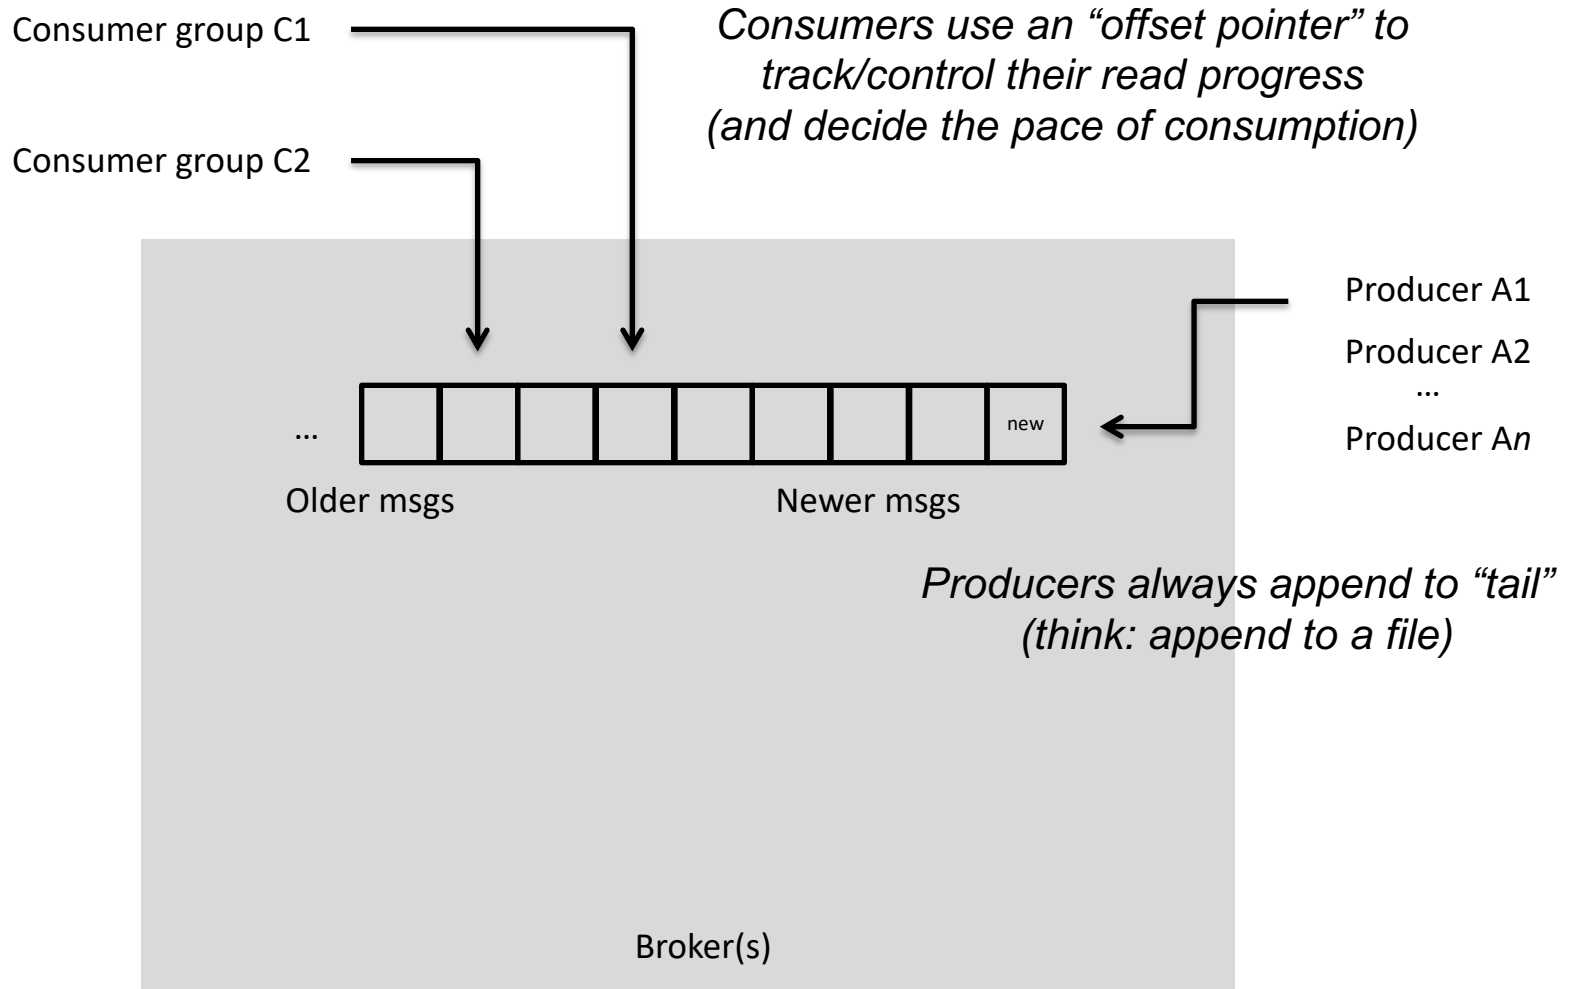
# Topics

- **Topic:** feed name to which messages are published
  - Example: “zerg.hydra”

*Kafka prunes “head” based on **age** or **max size** or “**key**”*




# Topics



# Partitions

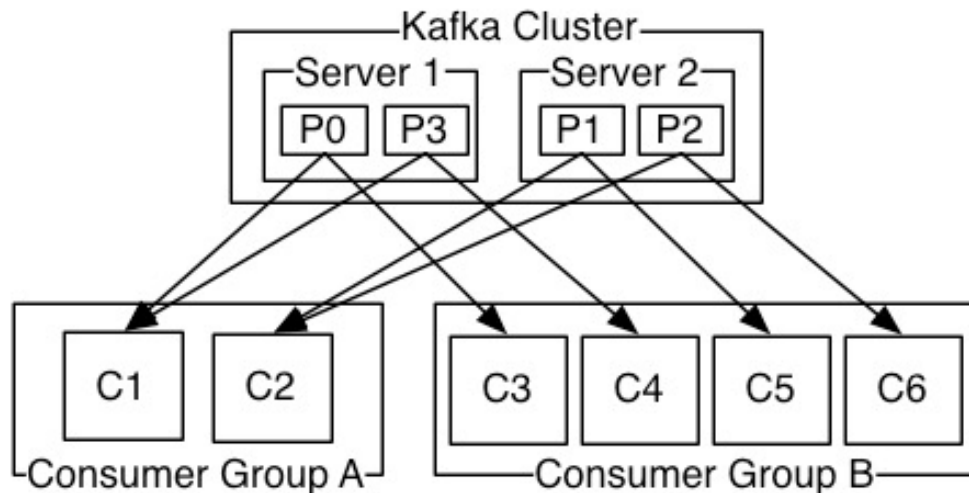
- A topic consists of **partitions**.
- Partition: **ordered + immutable** sequence of messages that is continually appended to

## Anatomy of a Topic



# Partitions

- #partitions of a topic is configurable
- #partitions determines **max** consumer (group) parallelism

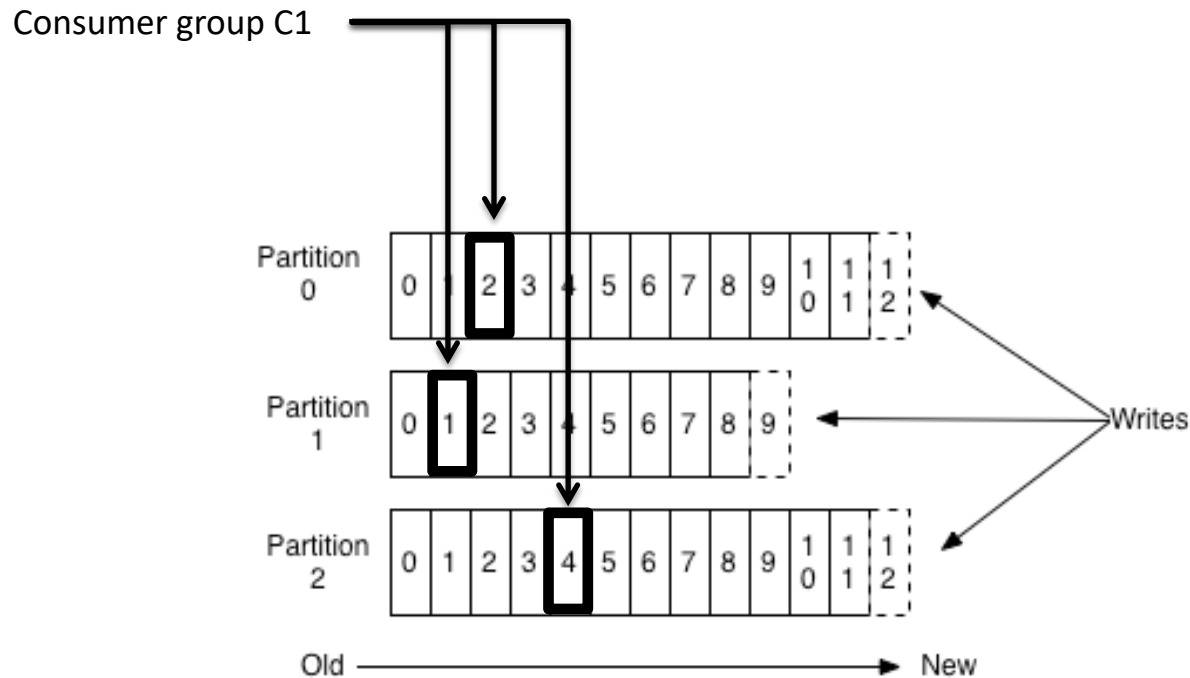


- Consumer group A, with 2 consumers, reads from a 4-partition topic
- Consumer group B, with 4 consumers, reads from the same topic



# Partition offsets

- **Offset:** messages in the partitions are each assigned a unique (per partition) and sequential id called the *offset*
  - Consumers track their pointers via (*offset, partition, topic*) tuples



# Replicas of a partition

- **Replicas:** “backups” of a partition
  - They exist solely to prevent data loss.
  - Replicas are never read from, never written to.
    - They do NOT help to increase producer or consumer parallelism!
  - Kafka tolerates  $(numReplicas - 1)$  dead brokers before losing data
    - LinkedIn: `numReplicas == 2` → 1 broker can die

# Brokers and Zookeeper

- Apache ZooKeeper is a high-reliability service for coordination of high-reliability distributed applications
- ZooKeeper
  - handles the leadership election of Kafka brokers
  - who the preferred leader node is for a given topic/partition pair
  - And other tasks to ensure system runs reliably despite failures of servers